



Erasmus Mundus joint
Master in Artificial Intelligence



Universitat
Pompeu Fabra
Barcelona



SAPIENZA
UNIVERSITÀ DI ROMA

Radboud Universiteit



UNIVERZA V LJUBLJANI
University of Ljubljana



Co-funded by
the European Union

Mitigating Objectness Bias and Region-to-Text Misalignment for Open-Vocabulary Panoptic Segmentation

Nikolay Kormushev

Supervisor: prof. dr. Matej Kristan

Co-supervisor: dr. Josip Šarić

May 2026



FRI

University of Ljubljana
Faculty of Computer and Information Science



Mitigating Objectness Bias and Region-to-Text Misalignment for Open-Vocabulary Panoptic Segmentation

Nikolay Kormushev

Erasmus Mundus Joint Master in Artificial Intelligence (EMA)

Faculty of Computer and Information Science

Supervisor: prof. dr. Matej Kristan
Cosupervisor: dr. Josip Šarić

Copyright: This work is licensed under a CC-BY-4.0 license. The software and results produced in this work are released under GPL-3.0-or-later at <https://github.com/nickormushev/OVRCOAT>.

Open-vocabulary panoptic segmentation remains hindered by two coupled issues: (i) mask selection bias, where objectness heads trained on closed vocabularies suppress masks of categories not observed in training, and (ii) limited regional understanding in vision-language models such as CLIP, which were optimized for global image classification rather than localized segmentation. We introduce OVRCOAT, a simple, modular framework that tackles both. First, a CLIP-conditioned objectness adjustment (COAT) updates background/foreground probabilities, preserving high-quality masks for out-of-vocabulary objects. Second, an open-vocabulary mask-to-text refinement (OVR) strengthens CLIP's region-level alignment to improve classification of both seen and unseen classes with markedly lower memory cost than prior fine-tuning schemes. The two components combine to jointly improve objectness estimation and mask recognition, yielding consistent panoptic gains. Despite its simplicity, OVRCOAT sets a new state of the art on ADE20K (+5.5% PQ) and delivers clear gains on Mapillary Vistas and Cityscapes (+7.1% and +3% PQ, respectively).

Panoptic segmentation | Open vocabulary | Computer vision

Panoptic segmentation provides a semantic visual scene decomposition by assigning a class-label to every pixel, while distinguishing individual object instances (things) from background regions (stuff) [1]. It thus unifies semantic and instance segmentation, combining class-level labelling with object-level differentiation, and delivers a comprehensive representation of objects and their surroundings.

Conventional panoptic segmentation models [2–10] are constrained by a fixed label set (vocabulary) of training data, thus limited only to the categories observed in training. This closed-vocabulary restriction poses a major obstacle for real-world deployment [11–15], where visual diversity exceeds the scope of curated datasets, motivating the development of generalisable segmentation paradigms.

One such paradigm is open-vocabulary panoptic segmentation, which leverages vision-language models (VLMs), for example, CLIP [16] and ALIGN [17]. These methods overcome the aforementioned limitations by first generating individual mask proposals, removing those with low objectness score, and classifying the rest by a VLM-based model. However, this introduces a mask selection bias, since the module for estimating the objectness score is trained on limited data. A recent study [18] noted that in practice, networks often provide very low scores on unlabelled or out-of-training-vocabulary masks, thus classifying them to the background. 1 visualises this bias, where state-of-the-art methods [12, 14] trained on datasets that did not include a *painting* class, incorrectly discard the corresponding masks.

Another challenge is that VLMs, such as CLIP [16] and ALIGN [17], are trained for global image classification rather than image regions localized by segmentation. While early approaches [12, 13] relied on frozen CLIP backbones, and proposed work-arounds, recent works such as OVSeg [19] and MAFT+ [14] propose involved fine-tuning strategies and achieve excellent open vocabulary semantic segmentation. However, in panoptic segmentation, only marginal improvements are observed on datasets like ADE20K [20] and notable drops on others.

To address the aforementioned challenges, we propose OVRCOAT a novel framework for open-vocabulary panoptic segmentation. The framework is built upon two key components. A CLIP-conditioned objectness adjustment (COAT) leverages CLIP [16] to refine mask selection by distinguishing foreground from background regions, effectively mitigating the mask selection bias present in previous models and improving the overall robustness. Open vocabulary mask-to-text feature refinement (OVR) introduces an objective that enhances CLIP's regional understanding and improves mask classification accuracy across both seen and unseen classes, while being considerably more memory-efficient than existing approaches. When combined, these two components form OVRCOAT, which jointly improves objectness estimation and mask classification, leading to improved panoptic generalization. Although fine-tuned specifically for panoptic segmentation, OVRCOAT retains strong generalization across other segmentation tasks.

Redukcija objektne pristranskosti ter neporavnosti med regijami in besedilom za panoptično segmentacijo z odprtim besednjakom

Panoptična segmentacija z odprtim besednjakom ostaja omejena zaradi dveh povezanih težav: (i) pristranskosti pri izbiri mask, kjer glave za oceno objektovskosti, naučene na zaprtih besednjakih, potlačijo maske kategorij, ki med učenjem niso bile opažene, in (ii) omejenega regionalnega razumevanja v vidno-jezikovnih modelih, kot je CLIP, ki so bili optimizirani za globalno klasifikacijo slik in ne za lokalizirano segmentacijo. Predstavimo OVRCOAT, preprost in modularen okvir, ki naslavlja oba problema. Najprej prilagoditev objektovskosti, pogojena s CLIP-om (COAT), posodobi verjetnosti ozadje/ospredje ter tako ohrani kakovostne maske za objekte zunaj besednjaka učenja. Nato izboljšava iz mask v besedilo za odprti besednjak (OVR) okrepi poravnavo CLIP-a na ravni regij, s čimer izboljša klasifikacijo tako vidnih kot nevidnih razredov ob bistveno nižjih pomnilniških stroških kot pri prejšnjih shemah finega prilagajanja. Obe komponenti se združita v skupno izboljšanje ocene objektovskosti in prepoznavanja mask, kar prinaša dosledne izboljšave panoptične segmentacije. Kljub svoji preprostosti OVRCOAT postavlja nov najboljši rezultat na ADE20K (+5.5% PQ) ter dosega jasne izboljšave na Mapillary Vistas in Cityscapes (ustrezno +7.1% in +3% PQ).

Panoptična segmentacija | Odprti besednjak | Računalniški vid



Figure 1. State-of-the-art (e.g., MAFT+ [14], FC-CLIP [12]) incorrectly discard detections on categories not present at training time (eg., paintings), leading to failed detection. OVRCOAT correctly classifies these instances by leveraging a CLIP-based update to background probabilities, improving performance on unseen classes.

In summary, our main contributions are:

- CLIP-conditioned objectness adjustment (COAT) mitigates the mask selection bias and improves previously unseen objects detection.
- Open vocabulary mask-to-text refinement (OVR) enhances CLIP’s localization and classification capabilities.

Compared to the current state-of-the-art [14], OVRCOAT proposes a substantially simpler architecture, which brings several benefits. It reduces the memory requirements, future extensions are simpler, and the proposed modules are easily deployable in related methods. OVRCOAT sets a new state-of-the-art on ADE20K [20], outperforming prior open-vocabulary panoptic works by +5.5% PQ, and delivers +7.1% PQ gains on Mapillary Vistas [21] and 3% PQ on Cityscapes [22], underscoring its generalisation across different domains.

1. Related Work

1.1. Vision-language models. Vision-language models (VLMs) learn mappings of images and text into a shared embedding space by training on large-scale collections of image-caption pairs. Seminal works such as CLIP [16] and ALIGN [17] adopt a contrastive objective that jointly optimizes image and text encoders to produce close embeddings for positive pairs while maximizing the distance for negative pairs. Such an objective is primarily retrieval-orientated, but also aligns well with classification tasks. These models demonstrate strong zero-shot performance by matching image features with text embeddings of class names, requiring no additional fine-tuning.

However, many downstream tasks, such as semantic and instance segmentation, require pixel-level recognition. Yet the contrastive objective enforces only global image-text alignment, resulting in representations with limited localization capabilities, which are crucial for dense prediction [23]. As a result, the naive application of VLMs to pixel-level tasks yields poor segmentation performance.

Various approaches have been considered to address these limitations. Some studies introduce additional pre-training objectives such as masked image reconstruction [24, 25], local-to-global view consistency [25], and region-to-text alignment [26]. In parallel, other studies repurpose pre-trained VLMs for dense prediction through post-training adaptation without relying on dense labels. For instance, these methods improve localization by refining attention mechanisms [27, 28] or by leveraging complementary representations from other foundation models [29–32]. Despite significant progress, training-free methods still achieve limited performance and most often lack the ability to distinguish individual object instances. To overcome these limitations, training-based open-vocabulary segmentation methods leverage dense labels to improve localization while retaining the generalization ability of vision–language pretraining.

1.2. Training-based open-vocabulary segmentation. While dense semantic labels enable direct fine-tuning of the VLM’s visual encoder, it is beneficial to employ a dedicated decoder to overcome the low output resolution and produce high-resolution segmentation maps. A line of work uses per-pixel decoders that operate on image–text cost volumes and

learn to refine the VLM’s predictions using dense supervision [33, 34]. While effective for semantic segmentation, extending these approaches to instance or panoptic segmentation remains highly challenging [35].

Consequently, most methods adopt mask-transformer decoders, which have proven to be effective architectures for all segmentation tasks [3, 4]. Early approaches [15, 19] employ mask transformers to generate mask proposals that are then cropped and classified using CLIP. However, this strategy requires multiple evaluations of the CLIP backbone, making it computationally expensive. More efficient variants [13, 36] mitigate this overhead by extracting image features once and reusing them to classify masks through a mask-pooling operator. FC-CLIP [12] further streamlines this design by employing a single frozen convolutional CLIP backbone together with a combined mask-transformer and CLIP-based classifier. MAFT+ [14] extends this idea by fine-tuning the backbone for mask classification, while preserving the vision–language alignment of CLIP. Despite these advances, mask-transformer-based methods still face two key challenges: suboptimal classification performance and missed detections of valid segments caused by biases in the training data [18]. Our work tackles these challenges directly.

2. OVRCOAT

We propose OVRCOAT, which features a clean architecture (2), with two complementary components. A CLIP-conditioned objectness adjustment (COAT) leverages CLIP’s classification confidence to debias the objectness score (Section 2.1), while an open-vocabulary mask-to-text refinement (OVR) adjusts CLIP’s regional understanding and improves mask-level classification accuracy across both seen and unseen categories (Section 2.2). The following subsections describe these two components in detail.




Figure 2. OVRCOAT extracts mask-to-text aligned features by CLIP image encoder and mask transformer, adjusted by a novel open-vocabulary mask-to-text refinement finetuning protocol (OVR). A new CLIP-conditioned objectness adjustment (COAT) module reduces the objectness bias in the mask transformer, improving detection of vocabulary instances not observed in training. Our main contributed blocks are denoted by red.

2.1. CLIP-Conditioned Objectness Adjustment. Mask-transformers-based [3, 4] panoptic methods propose a number of potential masks per image, remove the ones with low objectness scores and classify the rest into vocabulary-defined classes [12, 14]. The central interaction between vocabulary classes and objectness score occurs at the definition of the probability density function

$$\mathbf{p}_{cls} = [\mathbf{p}_{ens} \cdot p_{obj}, 1 - p_{obj}], \quad (1)$$

where $\mathbf{p}_{ens}$ is the standard ensembled per-vocabulary-class probability density [12] and p_{obj} is the objectness score (usually defined via $p_{obj} = 1 - p_{void}$), acting as a gating function – a low p_{obj} leads to early mask rejection. We make an insight that this score is computed by the mask transformer as a (normalized) similarity of the mask-wide feature with a pre-trained void token. It is thus substantially biased towards the training-time vocabulary, in particular, if test-time vocabulary instances occur in unlabeled regions of training images, the void token takes form that makes their classification as background more probable (i.e., low objectness).

On the other hand, CLIP does not contain the inherit bias due to its large-scale training, and effectively treats all categories as *seen* classes. We thus propose the following CLIP-based test-time vocabulary-guided bias correction. For a considered proposal mask $\mathbf{M}_i \in \{0, 1\}^{H \times W}$, we extract a feature representation $\mathbf{F}_{seg,i}$ by a global

mask pooling, i.e.,

$$\mathbf{F}_{\text{seg},i} = \frac{\sum_{u,v}^{H,W} \mathbf{M}_i(u,v) \cdot \mathbf{F}_{\text{img}}(u,v)}{\sum_{u,v}^{H,W} \mathbf{M}_i(u,v)}. \quad (2)$$

A probability distribution that the mask contains any of the test-time vocabulary categories is computed by softmax over dot-products between the mask feature and CLIP-encoded vocabulary terms embeddings $\mathbf{F}_{\text{txt}} \in \mathbb{R}^{N_{\text{cls}} \times E}$,

$$\mathbf{p}_{\text{CLIP},i} = \text{softmax}(\mathbf{F}_{\text{seg},i} \mathbf{F}_{\text{txt}}^\top) \in [0, 1]^{1 \times N_{\text{cls}}}. \quad (3)$$

If the mask feature agrees well with one of the classes, this results in a high corresponding probability. Alternatively, if it does not agree with any in particular, the distribution is driven towards uniform, effectively reducing the maximal probability. We thus estimate the CLIP-based classification certainty as $p_{\text{cer}} = \max_i \mathbf{p}_{\text{CLIP},i}$. Thus the objectness score in (1) is adjusted to

$$p'_{\text{obj}} = 1 - (1 - \gamma p_{\text{cer}})(1 - p_{\text{obj}}), \quad (4)$$

where p_{obj} is the initial mask-transformer objectness score and γ is the CLIP trust factor. Note that $\mathbf{p}_{\text{cer}}$ sets the upper bound on the adjusted objectness p'_{obj} , with higher values tightening the bound. To account for poor locality in CLIP embeddings, leading to noisy similarities, we set the trust factor to $\gamma = 0.5$.

2.2. Mask to Text Alignment. Since CLIP is trained for global image-to-text matching [16], it's per-pixel features do not align well with object boundaries, leading to inferior mask-level encoding [37]. To address this, we propose open-vocabulary mask-to-text refinement finetuning protocol (OVR) as shown in 3.




Figure 3. OVR optimizes a combined objective: a mask proposal accuracy $\mathcal{L}_{\text{M2F}}$ and mask classification $\mathcal{L}_{\text{cls}}$ for aligning per-pixel and text embeddings, while preserving vocabulary generalization.

In finetuning, embeddings are extracted by local pooling CLIP image encoder features for each mask proposal generated by the mask transformer. The features are matched to the CLIP embeddings of vocabulary terms and softmaxed, producing the class-wise probability density. Classification is supervised by a cross-entropy loss $\mathcal{L}_{\text{cls}}$ between the obtained class-wise density and the ground truth classification obtained by Hungarian matching [38] of the proposed masks with ground truth masks. In addition, to supervise mask transformer learning, the classical mask transformer loss [4] $\mathcal{L}_{\text{M2F}}$ is applied, leading to the following compounded loss

$$\mathcal{L} = \alpha \mathcal{L}_{\text{cls}} + \mathcal{L}_{\text{M2F}}. \quad (5)$$

Due to limited panoptic training datasets, care has to be taken in the refinement protocol design. Namely, the limited training-time visual examples and restricted class vocabulary pose overfitting risks, potentially weakening the vision-language alignment generalization. A two-stage approach is thus proposed. In the first stage, the CLIP image encoder (3) is kept frozen to pre-train the mask generation, without disrupting the CLIP image embedding space. In the second stage, the CLIP image encoder is unfrozen and training continues to jointly tune the embedding-mask alignment and mask proposal generation accuracy. As in [14] the final projection multilayer perceptron and normalization layers in the CLIP image encoder are frozen for restricted training.

Table 1. Analysis on standard datasets. Gray indicates a test set with same vocabulary as used in training.

Method	ADE20k [20]			Mapillary Vistas [21]			Cityscapes [22]			COCO [39]		
	PQ	SQ	RQ	PQ	SQ	RQ	PQ	SQ	RQ	PQ	SQ	RQ
MaskCLIP [15]	15.1	70.5	19.2	-	-	-	-	-	-	-	-	-
FreeSeg [40]	16.3	71.8	21.6	-	-	-	-	-	-	-	-	-
OPSNet [41]	19.0	52.4	23.0	-	-	-	41.5	67.5	50.0	52.4	83.5	62.1
ODISE [13]	23.4	78.1	28.3	14.2	61.0	17.2	23.9	75.3	29.0	55.4	-	-
FC-CLIP [12]	26.8	71.2	32.3	18.3	56.0	23.1	44.0	75.4	53.6	54.4	-	-
MAFT+pan [14]	27.1	73.5	32.9	15.7	55.5	19.8	38.3	70.2	46.9	50.3	82.2	60.3
OVRCOAT (Ours)	28.6 (+1.5)	77.3 (-0.8)	34.7 (+1.8)	19.6 (+1.3)	65.7 (+4.7)	24.8 (+1.7)	45.3 (+1.3)	78.7 (+3.3)	55.6 (+2.0)	54.6	82.9	65.1

3. Experiments

OVRCOAT uses the standard backbone, i.e., ConvNeXt-Large [12, 16, 42] from OpenCLIP [43], pretrained on LAION-2B [44]. The mask transformer architecture follows standard Mask2Former [4]. We train OVRCOAT using the AdamW optimiser with a weight decay 0.05, using a learning rate 1×10^{-4} in the first-stage training, and 5×10^{-5} in the second stage. The classification loss $\mathcal{L}_{cls}$ weight in (5) is set to $\alpha = 0.1$, while we use the same parameters as in [4] to train the mask transformer. The trust factor in (4) is set to $\gamma = 0.5$. We employ three NVIDIA A100 GPUs (40 GB each), and a batch size of 9.

In all of our experiments, the COCO [39] closed-vocabulary panoptic segmentation dataset is used for training OVRCOAT. We then evaluate it under two setups: open-vocabulary panoptic segmentation (3.1) and semantic segmentation (3.2.2).

3.1. Open-Vocabulary Panoptic Segmentation. OVRCOAT is compared with several state-of-the-art methods, i.e., MaskCLIP [15], FreeSeg [40], OPSNet [41], ODISE [13], FC-CLIP [12] and the panoptic version of the recent state-of-the-art MAFT+ [14] (MAFT+_{pan}) on three challenging out-of-vocabulary (OOV) datasets ADE20K [20], Cityscapes [22] and Mapillary Vistas [21]. Standard performance measures [1] are used: Panoptic Quality (PQ), Segmentation Quality (SQ) and Recognition Quality (RQ).

Results are reported in 1. OVRCOAT consistently outperforms previous methods across out-of-vocabulary datasets, on average outperforming the state-of-the-art MAFT+ in PQ by approximately 16%. Only on COCO [39], OVRCOAT slightly lags behind ODISE [13], with a 1.4% drop. This can be attributed to the COAT module, which is designed to improve generalisation by leveraging CLIP [16] for mask selection. While it reduces training biases in unseen classes, replacing the original specialised mask evaluation, which was trained on the COCO dataset, can hinder performance in those seen classes (we verify this empirically in 3.2.1). Nevertheless, despite this effect, OVRCOAT substantially outperforms ODISE on all other datasets, by up to 90% on Cityscapes [22].

OVRCOAT achieves substantial gains in the component metrics. In particular, on Mapillary Vistas [21], a remarkable 18% SQ performance improvement is observed compared to the current state-of-the-art MAFT+ [14] and 7.7% compared to ODISE [13]. An approximate 5.5% RQ improvement compared to top per-dataset performer is observed across all datasets. These results demonstrate the strong cross-domain generalisation of OVRCOAT. The generalisation result is particularly strong in the context of comparatively large MAFT+ PQ degradation outside of ADE20k [20], highlighting the challenges of fine-tuning.

3.2. Ablation study. We analyse the individual contributions of the proposed CLIP-conditioned objectness adjustment (COAT) and open-vocabulary mask-to-text refinement (OVR) to panoptic segmentation performance. The results are reported in 2.

We first isolate the effect of the proposed COAT by applying it to our baseline (FC-CLIP [12]). This modification alone consistently improves PQ across all out-of-vocabulary datasets, demonstrating gains ranging from 1.4% to 3%. We next evaluate the contribution of the mask-to-text refinement alone. The performance drops by 3.6% in PQ compared to original OVRCOAT, but improves by 1% – 5% compared to the baseline. Our training strategy, applied on the COCO dataset, is effective not only in enhancing performance on its validation set, encouraging CLIP to focus on the dataset’s distribution, but also in preserving a strong generalisation to unseen domains. Combining COAT and OVR, producing OVRCOAT, delivers clear advantages over the baseline, with improvements of up to 7% PQ, validating the complementary contributions of the two components, while incurring a small trade-off in performance on seen classes. Despite this, OVRCOAT still surpasses most state-of-the-art models (1).




Figure 4. Per-class PQ differences between OVRCOAT and FC-CLIP [12] on ADE20k [20] for classes with largest PQ differences. Mean values are reported weighted with ground truth class frequency (weighted), and treating all classes equally (unweighted). OVRCOAT yields substantial improvements on unseen categories, demonstrating stronger generalisation, with stable performance on seen classes.

Table 2. Ablation study of OVRCOAT individual modules.

COAT	OVR	ADE20k	Mapillary	Cityscapes	COCO
✗	✗	26.8	18.3	44.0	54.4
✓	✗	27.6	18.8	44.6	53.7
✗	✓	27.6	19.2	44.5	55.5
✓	✓	28.6	19.6	45.3	54.6

3.2.1. Per-class performance analysis. To more precisely disentangle the improvements introduced by OVRCOAT, we additionally analyze per-class panoptic quality (PQ) on ADE20K. We stratify categories into seen and unseen groups depending on whether they appear in the training set. 4 shows per-class PQ differences between the baseline and OVRCOAT for the 10 seen and 10 unseen classes with the largest absolute change, measured as $|PQ_{FC-CLIP}(c) - PQ_{OVRCOAT}(c)|$ for each class c . Additionally, we report the baseline enhanced with OVR, allowing us to isolate the specific contribution of COAT beyond this component.

For classes seen during training, OVRCOAT exhibits a mixed effect: categories less common in the COCO training set (≈ 5000 masks on average), such as *apparel*, *boat*, and *counter*, show clear gains, while those more frequently observed during training ($\approx 14\,000$ masks on average) experience only a moderate decrease. On average, panoptic quality across seen classes is minimally reduced for -0.05 pp, indicating that performance is largely preserved, while a +3.9 pp improvement is observed on unseen classes.

Conversely, unseen classes gain substantially from the combined effect of our mask-to-text refinement training strategy and the COAT module. For example, a dramatic relative improvement of 192% is observed for category *paintings*, which is commonly misclassified as background by the current state-of-the-art [45]. Overall, average relative improvement of 25% is observed across unseen classes. These results clearly highlight that OVRCOAT enhances generalisation to novel classes while simultaneously maintaining competitive accuracy on those seen during training.

3.2.2. Semantic Segmentation. For further insights, we evaluate the impact of our contributions on the semantic segmentation performance. 3 presents the results in terms of standard mean Intersection over Union (mIoU) on the PASCAL Context [46] and ADE20K [20] datasets. The first section presents the results of mask-transformer-based approaches from the literature. The second section incrementally presents the impact of our contributions OVR and COAT.

We first establish the notion of slight disadvantage of panoptic models compared to semantic segmentation models when evaluated on the semantic segmentation task. The results in 3 show that the current state-of-the-art MAFT+ variant trained specifically for semantic segmentation (MAFT+_{sem}) achieves the highest overall mIoU. It consistently outperforms the variant trained with panoptic labels (MAFT+_{pan}) by 1–9% in relative terms. This suggests that the additional demands of panoptic recognition consume part of the model’s capacity, thereby degrading the semantic segmentation performance.

Our open-vocabulary mask-to-text refinement (OVR) improves semantic segmentation performance over the baseline (FC-CLIP) in 4 out of 5 datasets, despite originally designed for panoptic segmentation. On the other hand, further inclusion of COAT slightly degrades the semantic segmentation performance. This behavior stems from the way mask transformers perform semantic inference. Unlike in panoptic segmentation, no masks are rejected. Instead, all predicted masks contribute to the aggregated class

distribution at every pixel. As a result, masks that would otherwise be discarded in panoptic inference due to low objectness, still influence the final per-pixel predictions. Hence, adjusting objectness has little real effect and actually introduces noise due to CLIP’s weaker localization.

Table 3. Open-vocabulary semantic segmentation performance (mIoU) across multiple datasets, with best performance bolded.

Model	Backbone	A-847	A-150	PC-459	PC-59	PAS-20
OpenSeg [17]	ALIGN	8.8	28.6	12.2	48.2	72.2
OVSeg [19]	ViT-L	9.0	29.6	12.4	55.7	94.5
SAN [47]	ViT-L	12.4	32.1	15.7	57.7	94.6
ODISE [13]	ViT-L	11.1	29.9	14.5	57.3	-
FC-CLIP* [12]	ConvNeXt-L	14.8	34.0	18.2	58.4	95.4
MAFT _{sem} [14]	ConvNeXt-L	15.1	36.1	21.6	59.4	96.5
MAFT _{pan} [14]	ConvNeXt-L	14.2	33.8	19.8	58.0	95.6
OVR	ConvNeXt-L	14.3	34.3	19.1	58.5	95.5
OVRCOAT	ConvNeXt-L	14.0	33.7	18.1	55.4	94.1

3.3. Architecture Design Validation Study.

3.3.1. Mask Classifier Validation. Since OVRCOAT fine-tunes the CLIP [16] image encoder on COCO [39], it is natural to investigate whether using only CLIP suffices for mask classification. 4 evaluates four different setups: the standard ensemble of mask-transformer and our CLIP_{OVR} predictions [12] with and without COAT, and using the fine-tuned CLIP_{OVR} alone, again with and without COAT. The results show that the ensemble with COAT (as used in OVRCOAT) outperforms all other variants. In particular, leveraging the mask-transformer probabilities as a stronger specialist on seen classes yields around a 14% relative improvement over using CLIP alone in both configurations. We also observe that COAT consistently boosts performance for both the ensemble and the standalone CLIP-based mask classifier.

Table 4. Comparison of various mask classification methods.

Mask Classifier	COAT	PQ	SQ	RQ
CLIP _{OVR}	✗	24.2	75.9	29.2
CLIP _{OVR}	✓	24.9	78.4	30.2
Ensemble	✗	27.6	74.1	33.4
Ensemble	✓	28.6	77.3	34.7

3.3.2. Classification of Perfect Masks. To further evaluate the effectiveness of our OVR fine-tuning approach and its impact on seen and unseen classes, we analyse the performance using a segmentation oracle. The oracle generates perfect binary masks, one for each ground truth panoptic segment, replacing those produced by the mask transformer. During inference, we follow the OVRCOAT pipeline, but completely omit the mask transformer, relying solely on the CLIP for classification.

We compare two variants of CLIP [16] using this oracle: the pre-trained OpenCLIP ConvNeXt-Large [42, 43], referred to as CLIP, and our OVR fine-tuned version (CLIP_{OVR}). Results in 5 show an 11% improvement on PQ in favour of CLIP_{OVR}. We also observe substantial gains of 9.2% on seen classes and 13% on unseen classes, indicating that our fine-tuning not only enhances performance on trained categories, but also generalises effectively. These results indicate that OVR effectively adapts CLIP to open-vocabulary panoptic segmentation.

Table 5. Evaluation of OVR using oracle (ground truth) masks.

Model	PQ	PQ _{seen}	PQ _{unseen}
CLIP	41.8	53.3	33.3
CLIP _{OVR}	46.4	58.2	37.6

3.3.3. Training strategy. Inspired by MAFT+ [14], OVR adopts a freezing strategy for the CLIP ConvNeXt-L backbone [16, 42], targeting the final MLP and normalisation layers, but discarding their representation consistency (RC) loss. To support this choice, we

compare OVR against five variants that employ different strategies to preserve CLIP’s embedding alignment. For a controlled comparison, we omit COAT in all variants of this experiment.

We define the variants of the training loss as $\mathcal{L}_c \in \mathcal{L}, RC, G$, where $\mathcal{L}$ denotes the base loss introduced in (5). The variant $RC = \mathcal{L} + \mathcal{L}_{RC}$ incorporates the MAFT+ [14] RC loss, while $G = \mathcal{L} + \mathcal{L}_{Gram}$ adds a Gram Matrix consistency loss inspired by DINOv3 [48]. Furthermore, we define the state of the CLIP [16] image encoder as $s \in F, U$, where F and U denote a frozen and unfrozen state, respectively. This allows us to denote each model variant as $OVR_{\mathcal{L}_c, s}$, with the baseline represented as $OVR_{B, F}$. The results are shown in 6.

Table 6. Comparison of finetuning strategies on ADE20K [20].

Method	PQ	SQ	RQ
$OVR_{B, U}$	27.2	72.4	33.1
$OVR_{RC, U}$	27.7	73.0	33.7
$OVR_{G, U}$	27.7	71.3	33.6
$OVR_{G, F}$	27.6	73.5	33.4
$OVR_{RC, F}$	27.7	71.2	33.6
OVR	27.6	73.6	33.4

Training without additional constraints ($OVR_{B, U}$) results in only a 1.5% decrease in PQ compared to OVR, while achieving performance comparable to MAFT+ [14] (cf. 1) with a marginal gain of 0.1 PQ (0.03%). This robustness is attributed to the stage-one training and the lower learning rate used during the second-stage tuning of CLIP [16]. Aggressively enforcing feature constraints via RC or Gram matrix losses [14, 15] does not have a significant impact on models performance with respect to OVR. Across all configurations, the best-performing variants achieve results within 1% of OVR’s performance. These findings justify freezing the final layers without additional auxiliary losses, providing a favourable trade-off between simplicity and stability while avoiding the complexity and sensitivity introduced by constraint-based training.

3.3.4. Memory Efficiency. We further analyse the memory usage of OVRCOAT. Results in 7 show that OVRCOAT requires substantially less memory compared to the state-of-the-art MAFT+ [14]. For a batch size of 1, the average memory usage is reduced by approximately 56% (16.1 GB). This is due to the streamlined design of our fine-tuning strategy OVR, which does not require content-dependent vocabulary embeddings or additional mask-overlap computations.

Considering that the reduction in memory consumption comes with an improvement in 5.5% PQ, OVRCOAT brings several practical benefits: it accelerates experimentation, enables use of larger batches (which is particularly important in modern unsupervised learning), and caters to the trend of exploring ever larger models, where memory consumption is a critical limiting factor. Finally, it also enables training on resource-constrained hardware, offering far greater accessibility.

3.3.5. CLIP Trust Factor. As a final validation, we evaluate the sensitivity of OVRCOAT to the CLIP trust weight γ used in (4). 5 reports the PQ values on ADE20K [20] with respect to different values of γ . As γ increases towards 0.5, the performance increases




Figure 5. OVRCOAT performance on ADE20K [20] remains stable over a range of trust factor γ values.




Figure 6. Qualitative analysis. OVRCOAT consistently recovers more object masks, including regions that MAFT+ [14] and FC-CLIP [12] misclassify as background, e.g., playthings in the bedroom and bottles in the bar. It also produces more accurate segmentations.

Table 7. Average training-time memory consumption.

Model	PQ (ADE20K)	Memory [GB/image]
MAFT+	27.1	27.0
OVRCOAT	28.6	12.5

to 28.6 PQ, then gradually decreases to 22.8 PQ (a 20% drop) at $\gamma = 1.0$. The graph suggests that OVRCOAT is stable across a range of CLIP trust values, with a broad range of values resulting in top PQ score.

3.4. Qualitative analysis. 6 provides qualitative comparison of OVRCOAT with FC-CLIP [12] and MAFT+ [14]. We observe that OVRCOAT consistently identifies more meaningful object masks than the state-of-the-art, supporting the generalization improvements reflected in previous experiments. In the *bedroom* and *children* scenes, OVRCOAT successfully recovers additional playthings and balls that competing models misclassify as background. In the *museum* example, OVRCOAT detects finer architectural details such as doors and pillars. Moreover, OVRCOAT produces finer, object-centric segmentations. For instance, it separates multiple bottles in the *bar* scene instead of a larger shelf mask as in MAFT+ [14]. This behaviour highlights a strengthened general object understanding, particularly benefiting the detection of smaller and unseen objects.

4. Conclusion

We addressed two pressing obstacles in open-vocabulary panoptic segmentation: the mask-selection bias induced by mask transformer training on a closed vocabulary dataset, and reduced accuracy of region-level understanding in vision-language models trained for global classification. We introduced OVRCOAT, a minimal and modular framework that tackles both: COAT adjusts background/foreground probabilities using CLIP to retain object masks for unseen categories, while OVR refines mask-to-text alignment to improve recognition of both seen and unseen classes.

OVRCOAT’s design brings practical benefits: clean integration into standard Mask2Former pipelines and substantially lower memory demands than recent fine-tuning approaches, while remaining robust across datasets. Empirically, OVRCOAT consistently surpasses prior open-vocabulary panoptic methods on challenging out-of-vocabulary benchmarks, and ablations verify the individual components. COAT reduces selection bias and boosts recall of novel objects, OVR strengthens regional alignment, and their combination yields the most reliable panoptic improvements. Analyses with oracle masks confirm the alignment gains, probability-fusion studies favor combining CLIP with mask-transformer scores, and training-strategy experiments show that strong results are attainable without complex representation-consistency losses.

Looking forward, we see OVRCOAT as a foundation for broader open-world perception, extending to video panoptics with temporal consistency, coupling with promptable or interactive segmentation. As stronger VLMs and data engines emerge, OVRCOAT’s modularity should make such upgrades straightforward, further closing the gap between closed-set training and open-world deployment.

ACKNOWLEDGEMENT. This work was supported in part by research programs and projects P2-0214, J2-60054 and SLAIF grant no. 101254461 and by ARNES/SLING supercomputer clusters.

Bibliography

1. Kirillov A, He K, Girshick R, Rother C, Dollár P (2019) Panoptic segmentation in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 9404–9413.
2. Carion N et al. (2020) End-to-end object detection with transformers in *European conference on computer vision*. (Springer), pp. 213–229.
3. Cheng B, Schwing A, Kirillov A (2021) Per-pixel classification is not all you need for semantic segmentation. *Advances in neural information processing systems* 34:17864–17875.
4. Cheng B, Misra I, Schwing AG, Kirillov A, Girdhar R (2022) Masked-attention mask transformer for universal image segmentation in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 1290–1299.
5. Wang H, Zhu Y, Adam H, Yuille A, Chen LC (2021) Max-deeplab: End-to-end panoptic segmentation with mask transformers in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 5463–5474.
6. Li Z et al. (2022) Panoptic segformer: Delving deeper into panoptic segmentation with transformers in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 1280–1289.
7. Yu Q et al. (2022) Cmt-deeplab: Clustering mask transformers for panoptic segmentation in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 2560–2570.
8. Jain J et al. (2023) Oneformer: One transformer to rule universal image segmentation in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 2989–2998.
9. Li F et al. (2023) Mask dino: Towards a unified transformer-based framework for object detection and segmentation in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 3041–3050.
10. Sun S et al. (2023) Remax: Relaxing for better training on efficient panoptic segmentation. *Advances in Neural Information Processing Systems* 36:73480–73496.
11. Zhao H, Puig X, Zhou B, Fidler S, Torralba A (2017) Open vocabulary scene parsing in *Proceedings of the IEEE International Conference on Computer Vision*. pp. 2002–2010.
12. Yu Q, He J, Deng X, Shen X, Chen LC (2023) Convolutions die hard: Open-vocabulary segmentation with single frozen convolutional clip. *Advances in Neural Information Processing Systems* 36:32215–32234.
13. Xu J et al. (2023) Open-vocabulary panoptic segmentation with text-to-image diffusion models in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 2955–2966.
14. Jiao S et al. (2024) Collaborative vision-text representation optimizing for open-vocabulary segmentation in *European Conference on Computer Vision*. (Springer), pp. 399–416.
15. Ding Z, Wang J, Tu Z (2022) Open-vocabulary universal image segmentation with maskclip. *arXiv preprint arXiv:2208.08984*.
16. Radford A et al. (2021) Learning transferable visual models from natural language supervision in *International conference on machine learning*. (PMLR), pp. 8748–8763.
17. Jia C et al. (2021) Scaling up visual and vision-language representation learning with noisy text supervision in *International conference on machine learning*. (PMLR), pp. 4904–4916.
18. Šarić J, Martinović I, Kristan M, Šegvić S (2025) What holds back open-vocabulary segmentation? in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV) Workshops*. pp. 4256–4266.
19. Liang F et al. (2023) Open-vocabulary semantic segmentation with mask-adapted clip in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 7061–7070.
20. Zhou B et al. (2017) Scene parsing through ade20k dataset in *Proceedings of the IEEE conference on computer vision and pattern recognition*. pp. 633–641.
21. Neuhold G, Ollmann T, Rota Buló S, Kotschieder P (2017) The mapillary vistas dataset for semantic understanding of street scenes in *Proceedings of the IEEE international conference on computer vision*. pp. 4990–4999.
22. Cordts M et al. (2016) The cityscapes dataset for semantic urban scene understanding in *Proceedings of the IEEE conference on computer vision and pattern recognition*. pp. 3213–3223.
23. Rao Y et al. (2022) Denseclip: Language-guided dense prediction with context-aware prompting in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 18082–18091.
24. Luo H, Bao J, Wu Y, He X, Li T (2023) Segclip: Patch aggregation with learnable centers for open-vocabulary semantic segmentation in *International Conference on Machine Learning*. (PMLR), pp. 23033–23044.
25. Tschannen M et al. (2025) Siglip 2: Multilingual vision-language encoders with improved semantic understanding, localization, and dense features. *arXiv preprint arXiv:2502.14786*.
26. Xu J et al. (2022) Groupvit: Semantic segmentation emerges from text supervision in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 18134–18144.
27. Zhou C, Loy CC, Dai B (2022) Extract free dense labels from clip in *European conference on computer vision*. (Springer), pp. 696–712.
28. Wang F, Mei J, Yuille A (2024) Sclip: Rethinking self-attention for dense vision-language inference in *European Conference on Computer Vision*. (Springer), pp. 315–332.
29. Wysoczańska M et al. (2024) Clip-dinoiser: Teaching clip a few dino tricks for open-vocabulary semantic segmentation in *European Conference on Computer Vision*. (Springer), pp. 320–337.
30. Zhang D, Liu F, Tang Q (2025) Corrcip: Reconstructing patch correlations in clip for open-vocabulary semantic segmentation in *Proceedings of the IEEE/CVF International Conference on Computer Vision*. pp. 24677–24687.
31. Wang J et al. (2025) Declip: Decoupled learning for open-vocabulary dense perception in *Proceedings of the Computer Vision and Pattern Recognition Conference*. pp. 14824–14834.
32. Lan M et al. (2024) Proxycip: Proxy attention improves clip for open-vocabulary segmentation in *European Conference on Computer Vision*. (Springer), pp. 70–88.
33. Cho S et al. (2024) Cat-seg: Cost aggregation for open-vocabulary semantic segmentation in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. pp. 4113–4123.
34. Xie B, Cao J, Xie J, Khan FS, Pang Y (2024) Sed: A simple encoder-decoder for open-vocabulary semantic segmentation in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 3426–3436.
35. Martinović I, Šarić J, Oršić M, Kristan M, Šegvić S (2025) Dearli: Decoupled enhancement of recognition and localization for semi-supervised panoptic segmentation. *arXiv preprint arXiv:2507.10118*.
36. Ghiasi G, Gu X, Cui Y, Lin TY (2022) Scaling open-vocabulary image segmentation with image-level labels in *European conference on computer vision*. (Springer), pp. 540–557.
37. Shi B et al. (2024) Umg-clip: A unified multi-granularity vision generalist for open-world understanding in *European Conference on Computer Vision*. (Springer), pp. 259–277.
38. Kuhn HW (1955) The hungarian method for the assignment problem. *Naval research logistics quarterly* 2(1-2):83–97.
39. Lin TY et al. (2014) Microsoft coco: Common objects in context in *European Conference on Computer Vision (ECCV)*.
40. Qin J et al. (2023) Freeseg: Unified, universal and open-vocabulary image segmentation in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. pp. 19446–19455.
41. Chen X, Li S, Lim SN, Torralba A, Zhao H (2023) Open-vocabulary panoptic segmentation with embedding modulation in *Proceedings of the IEEE/CVF International Conference on Computer Vision*. pp. 1141–1150.
42. Liu Z et al. (2022) A convnet for the 2020s in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 11976–11986.
43. Ilharco G et al. (2021) Openclip (v0.1) (<https://doi.org/10.5281/zenodo.7301701>). Zenodo.
44. Schuhmann C et al. (2022) Laion-5b: An open large-scale dataset for training next generation image-text models. *Advances in neural information processing systems* 35:25278–25294.
45. Šarić J, Martinović I, Kristan M, Šegvić S (2025) What holds back open-vocabulary segmentation? in *Proceedings of the IEEE/CVF International Conference on Computer Vision*. pp. 4256–4266.
46. Mottaghi R et al. (2014) The role of context for object detection and semantic segmentation in the wild in *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.
47. Xu M, Zhang Z, Wei F, Hu H, Bai X (2023) Side adapter network for open-vocabulary semantic segmentation in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. pp. 2945–2954.
48. Siméoni O et al. (2025) Dinov3. *arXiv preprint arXiv:2508.10104*.